

Daily Work Report

Sumukh Naidu H M | AI Engineer | Wednesday, September 09, 2026

Project	Corporate Intelligence Extraction POC
Pipeline	Discover → Crawl → Clean → Classify → Extract → Resolve → Deliver

Today's Work

- Built out the CWI-09 benchmark / evaluation harness for the extraction pipeline, structured as its own top-level **benchmark/** package (schema.py, normalize.py, merge.py, labels.py, scorer.py, report.py, cli.py).
- Harness reads only the existing **data/cleaned/*.json** output — zero network calls — and scores against hand-labeled ground truth stored per company in **data/labels/*.json**.
- Implemented CLI commands (**label** and **run**) for labeling ground truth and executing scoring runs.
- Implemented six-outcome scoring logic: correct, incorrect, missing, false_positive, correctly_abstained, and unlabeled — computing precision, coverage, and false-population-rate metrics, with safe null handling on zero-denominator cases.
- Added per-field match-type rules (exact / fuzzy / phone / URL) so comparisons account for how each field type should reasonably be matched.
- Added a naive, eval-only per-company merge step to stand in for the not-yet-built Resolve stage, plus explicit "pending" sentinels for citation_validity, since CWI-10 (citation verification) has not been built yet — avoided faking a number for an unmeasured metric.
- Wrote and passed all 43 unit tests covering the harness (test_labels, test_merge, test_normalize, test_scorer).

Upcoming Work

- Starting work on Odoo CRM integration based on a requirement document that has been provided.
- Will review Odoo's open-source codebase to identify which existing CRM features are actually relevant to the stated requirement, rather than building from scratch.
- Goal is to shortlist only the features genuinely needed for the requirement, avoiding unnecessary scope from Odoo's full CRM feature set.